

From Text to Insight: Disaster Detection with NLP

1. Problem Description

I participated in the Kaggle competition "Natural Language Processing with Disaster Tweets," which required me to build a machine learning model to classify tweets as either related to a real disaster or not:

Objective:

0: Non-disaster tweet.

1: Disaster tweet.

Challenges:

Tweets are short and may contain unstructured data, including typos, symbols, and emojis.

This task required natural language processing (NLP) techniques for feature extraction and classification.

2. Data Processing and Feature Extraction

Data Description

Training dataset (`train.csv`):

Contained 7613 records with tweet text (`text`), keywords (`keyword`), location (`location`), and labels (`target`).

Test dataset (`test.csv`):

Contained tweet text (`text`), keywords (`keyword`), and location (`location`), but without labels.

Data Cleaning

I filled missing values in the `keyword` and `location` columns with '`missing`' and '`unknown`', respectively.

I cleaned the text by removing URLs, punctuation, and extra spaces, and converted all text to lowercase.

Feature Extraction

I used the TF-IDF method to convert text into numerical features, selecting the top 5000 most frequent words for embedding.

3. Model Selection and Training

Model Selection

I chose Logistic Regression as the baseline model because it is simple, efficient, and interpretable.

Performance Results:

Accuracy: 80.4%.

The confusion matrix indicated that the model had lower recall when classifying disaster-related tweets (class 1).

Results Analysis

Classification Report:

vbnet

Copy code

Precision: 0.82 (Disaster Class)

Recall: 0.69 (Disaster Class)

F1-Score: 0.75 (Disaster Class)

Confusion Matrix:

The model misclassified some disaster-related tweets as non-disaster, leading to lower recall for class 1.

4. Submission and Testing Results

Test Data Prediction

I applied the same cleaning and feature extraction process to the test dataset.

Using the trained logistic regression model, I predicted the labels for the test data.

Submission File

I generated the `submission.csv` file in the required format:

bash

Copy code

`id,target`

`0,1`

`1,0`

`2,1`

After submitting to Kaggle, I reviewed my ranking on the leaderboard.

5. Learnings and Future Improvements

Key Takeaways

I gained experience in building a complete NLP pipeline, including data cleaning, feature extraction, model training, and evaluation.

This project helped me understand how to handle real-world text classification problems.

Future Improvements

I plan to experiment with deep learning models, such as LSTM, to capture sequential patterns in text.

I aim to explore more advanced embedding techniques like Word2Vec or GloVe to improve feature representation.

Fine-tuning hyperparameters could help improve recall and overall performance.

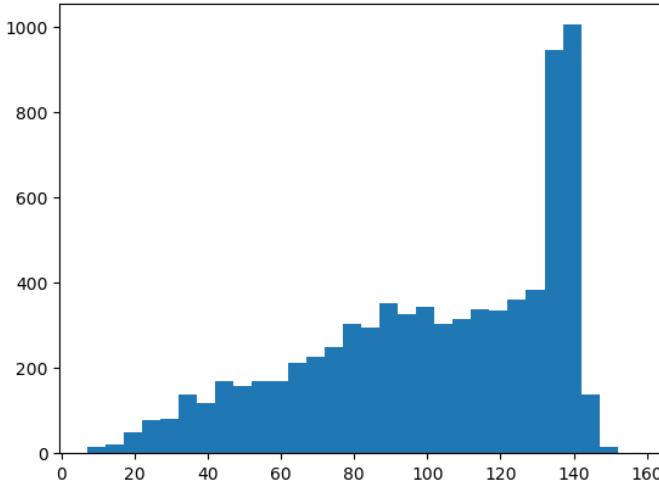
```
[19]: import pandas as pd  
data = pd.read_csv("train.csv")  
print(data.shape)  
data.head()
```

(7613, 5)

```
[19]:   id keyword location          text  target  
0    1      NaN     NaN  Our Deeds are the Reason of this #earthquake M...      1  
1    4      NaN     NaN  Forest fire near La Ronge Sask. Canada      1  
2    5      NaN     NaN  All residents asked to 'shelter in place' are ...      1  
3    6      NaN     NaN  13,000 people receive #wildfires evacuation or...      1  
4    7      NaN     NaN  Just got sent this photo from Ruby #Alaska as ...      1
```

```
[20]: import matplotlib.pyplot as plt  
data['text_length'] = data['text'].apply(len)  
plt.hist(data['text_length'], bins=30)  
plt.title("Tweet Length Distribution")  
plt.show()
```

Tweet Length Distribution



```
[21]: print(data.isnull().sum())

id          0
keyword     61
location    2533
text         0
target       0
text_length  0
dtype: int64

[22]: data['keyword'] = data['keyword'].fillna('missing')
data['location'] = data['location'].fillna('unknown')

print(data.isnull().sum())

keyword_counts = data['keyword'].value_counts()
keyword_counts[:10].plot(kind='bar', title='Top 10 Keywords')
plt.show()

location_counts = data['location'].value_counts()
location_counts[:10].plot(kind='bar', title='Top 10 Locations')
plt.show()

import re
def clean_text(text):
    text = text.lower()
    text = re.sub(r'http\S+', '', text)
    text = re.sub(r'[a-zA-Z\s]', '', text)
    text = re.sub(r'\s+', ' ', text).strip()
    return text

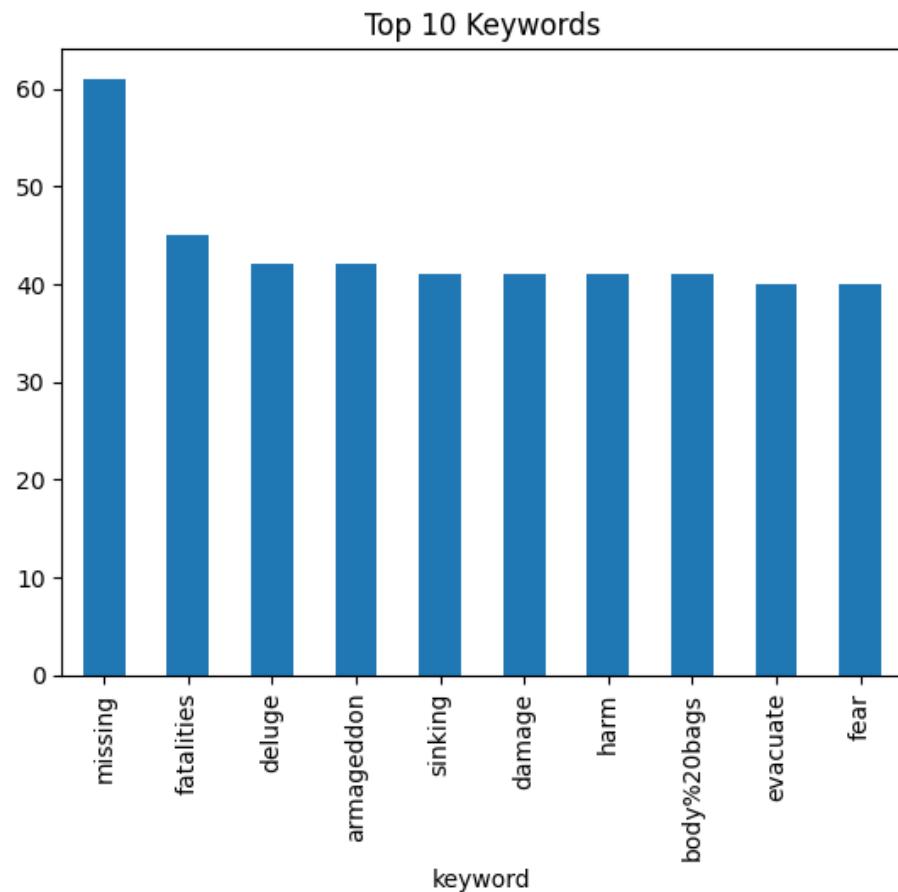
data['clean_text'] = data['text'].apply(clean_text)

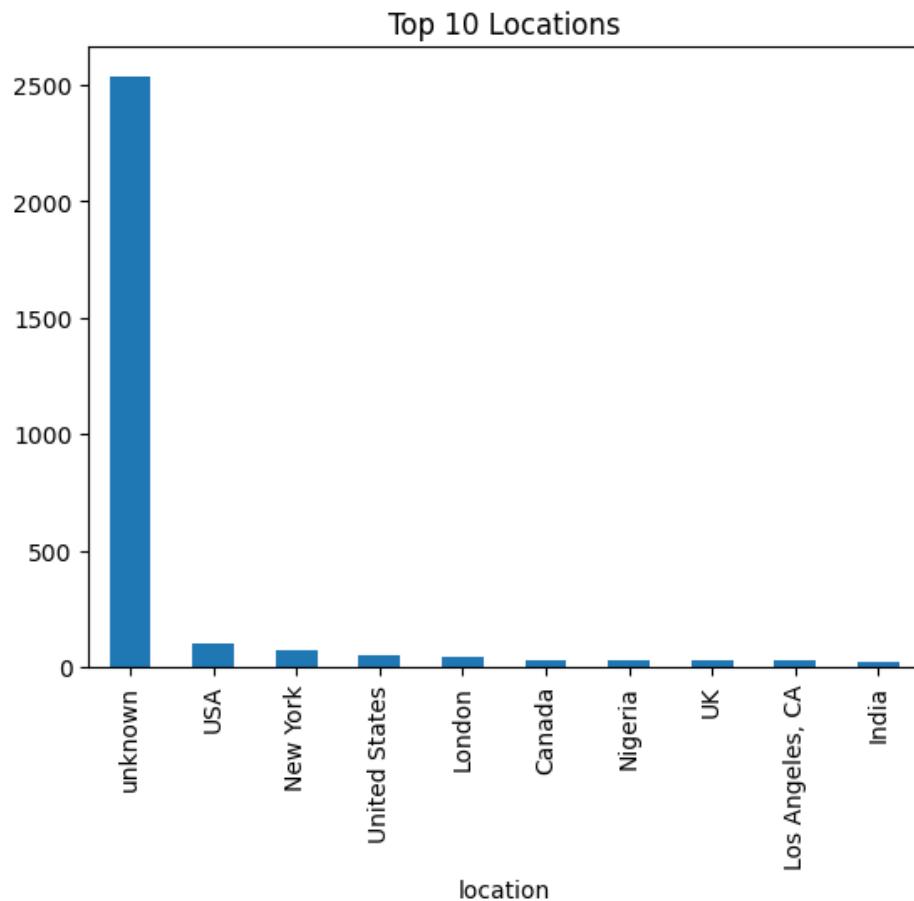
from sklearn.feature_extraction.text import TfidfVectorizer
tfidf = TfidfVectorizer(max_features=5000)
X = tfidf.fit_transform(data['clean_text']).toarray()
y = data['target']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print("Training data shape:", X_train.shape)
print("Test data shape:", X_test.shape)
```

```
id          0
keyword     0
location    0
text         0
target       0
text_length  0
dtype: int64
```





```
Training data shape: (6090, 5000)
Test data shape: (1523, 5000)
```

```
[29]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt

logistic_model = LogisticRegression(max_iter=1000, random_state=42)
logistic_model.fit(X_train, y_train)

y_pred = logistic_model.predict(X_test)

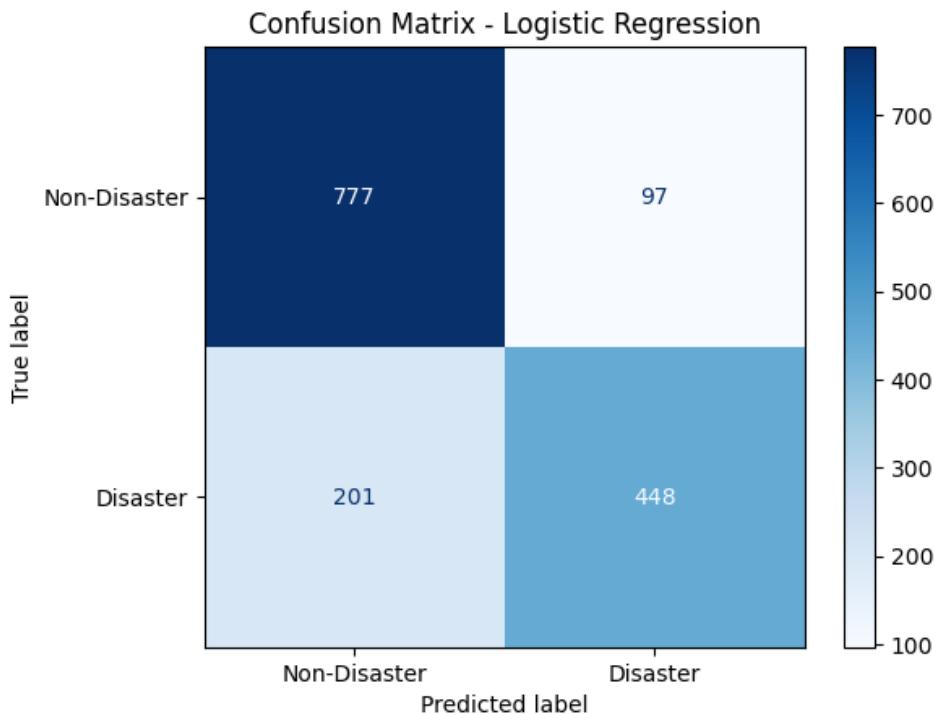
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

ConfusionMatrixDisplay.from_predictions(y_test, y_pred, display_labels=['Non-Disaster', 'Disaster'], cmap='Blues')
plt.title("Confusion Matrix - Logistic Regression")
plt.show()
```

```
Accuracy: 0.804333552199606
Classification Report:
      precision    recall  f1-score   support

          0       0.79      0.89      0.84      874
          1       0.82      0.69      0.75      649

   accuracy                           0.80      1523
  macro avg       0.81      0.79      0.79      1523
weighted avg       0.81      0.80      0.80      1523
```



```
[31]: test_data = pd.read_csv("test.csv")

test_data['keyword'] = test_data['keyword'].fillna('missing')
test_data['location'] = test_data['location'].fillna('unknown')
test_data['clean_text'] = test_data['text'].apply(clean_text)

X_test_final = tfidf.transform(test_data['clean_text']).toarray()

test_predictions = logistic_model.predict(X_test_final)

submission = pd.DataFrame({'id': test_data['id'], 'target': test_predictions})
submission.to_csv("submission.csv", index=False)

print("Submission file 'submission.csv' created successfully!")
```

Submission file 'submission.csv' created successfully!